

Introduction to Deep Learning in Vision: Basics, Optimization, Networks and Coding

Yue Zhang

Department of Mathematics, Applied Mathematics and Statistics
Case Western Reserve University

February 12, 2018

Outline

① Overview

② Convolutional Neural Networks Basics

Basic Operations

Optimization Problems in CNNs

③ Optimization

④ Networks Variants

⑤ Coding

Overview of the Journey

CNNs Basics

Standard Operations:

Convolutions , BN
Pooling,
Nonlinearities *etc.*

Special Terms:

Deconvolution,
Upsampling, Skip
Connection, Dense-
block *etc.*

Popular Losses:

Entropy based,
Standard l_2 and l_1 ,
Adversarial Losses,
etc.

Optimization

Mathematics:

Backpropagation on
CNNs

Popular first order methods:

SGD,
Momentum, Nesterov
Acceleration, Adam,
RMSprop *etc.*

Network Variants

LeNet (1998),
AlexNet (2012),
VGGNet (2014),
GoogLeNet (2014),
FCN (2014),
ResNet (2015),
U-Net (2015),
{ SegNet (2015),
DenseNet (2017),
Dense-UNet (2017) }
GAN (2014)
{ C-GAN (2014),
Cycle-GAN (2017) }

Coding

Frameworks:

Caffe (Berkeley),
Caffe2 (Facebook),
Theano (Bengio),
Torch (Facebook),
Pytorch (Facebook),
Tensorflow (Google)

Outline

- ① Overview
- ② Convolutional Neural Networks Basics
 - Basic Operations
 - Optimization Problems in CNNs
- ③ Optimization
- ④ Networks Variants
- ⑤ Coding


Basic Operations in CNNs

Standard operations in a Convolutional Neural Network:

- Convolution
- Pooling
- Batch Normalization
- Nonlinear Activation
- Others: Deconvolution, Upsampling, Skip Connections *etc.*

A Quick Example: UNet


One of the most popular networks in semantic segmentation.



Olaf Ronneberger, Philipp Fischer, Thomas Brox. *U-Net: Convolutional Networks for Biomedical Image Segmentation*, MICCAI 2015.

A Quick Example: UNet

One of the most popular networks in semantic segmentation.



Olaf Ronneberger, Philipp Fischer, Thomas Brox. *U-Net: Convolutional Networks for Biomedical Image Segmentation*, MICCAI 2015.

Convolution and Convolution Layer

(Conv + BN + ReLU)

Keywords in Convolution: • **Kernel Size** • **Stride** • **Padding**

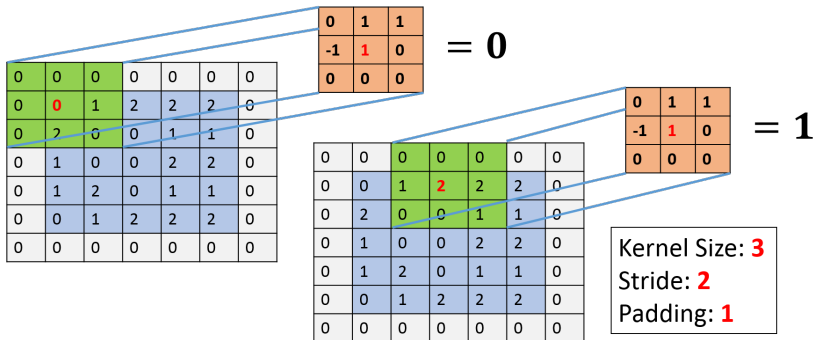


Figure: Illustration of Convolution¹.

¹To make it simple, the kernel is already **rotated**. Only point-wise product and summation is needed

Convolutions in CNNs: (Conv + BN + ReLU)

Different types of convolutions.

- **Convolution:** (with/without) padding, $(1/> 1)$ stride.
- **Transposed Convolution (Deconvolution):** (with/without) padding, $(1/> 1)$ stride.
- **Dilated Convolution.**

See the attached HTML file.

Helpful reading: [Vincent Dumoulin, Francesco Visin](#). *A guide to convolution arithmetic for deep learning*.

Convolutional Layer (1st): (Conv + BN + ReLU)

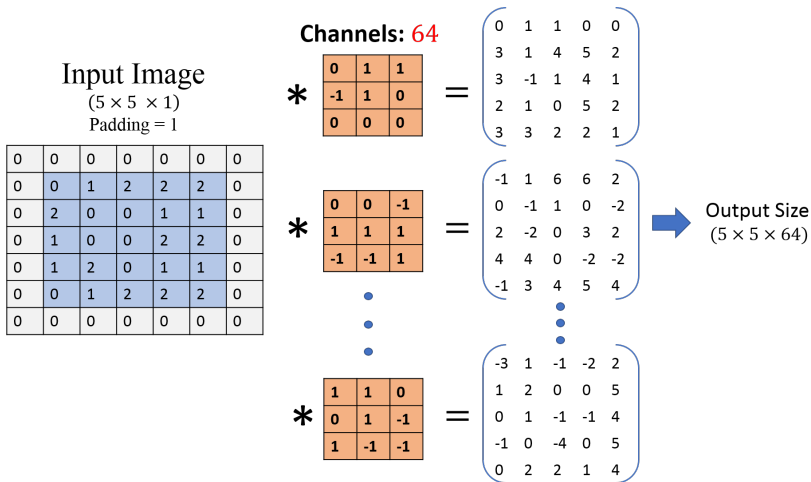


Figure: First Layer of CNN².

²Again, all the kernels are already **rotated**. Only point-wise product and summation is needed.

Convolutional Layer (2nd): (Conv + BN + ReLU)




Figure: Each channel has 64 different 3×3 filters. Each filter convolves with only *one* channel of the input feature map!

A Short Break: A few questions ...

Suppose we have $4 \times 512 \times 1$ image as network input. That is, (batch size) 4 images where each of them is $512 \times 512 \times 1$ (gray images).

Then,

- how many parameters (numbers in filters) do we have so far for the first two layers?

A Short Break: A few questions ...

Suppose we have $4 \times 512 \times 1$ image as network input. That is, (batch size) 4 images where each of them is $512 \times 512 \times 1$ (gray images).

Then,

- how many parameters (numbers in filters) do we have so far for the first two layers?
 - $64 \times 3 \times 3 + 128 \times 64 \times 3 \times 3 = 576 + 73728 = 74,304$

A Short Break: A few questions ...

Suppose we have $4 \times 512 \times 1$ image as network input. That is, (batch size) 4 images where each of them is $512 \times 512 \times 1$ (gray images).

Then,

- how many parameters (numbers in filters) do we have so far for the first two layers?
 - $64 \times 3 \times 3 + 128 \times 64 \times 3 \times 3 = 576 + 73728 = 74,304$
- since both the batches, feature maps are stored in memory, how much memory do we need? (suppose padding = 1, stride = 1)

$$\begin{aligned} & 4 \times 512 \times 512 \times 1 + 4 \times 512 \times 512 \times 64 + 4 \times 512 \times 512 \times 128 \\ & \approx 1M + 67M + 134M = 202M \\ & = 202M \times 4 \text{ bytes} \approx 770\text{MB} \quad (202M \times 4/1024^2) \end{aligned}$$

Some background. Almost all deep learning models are trained on GPUs. A typical GPU now has $6 \sim 8$ GB memory. Advance GPUs has 12 GB memory (*e.g.* Nvidia GeForce GTX TITAN Z $\sim \$1.5K$ on amazon).

Batch Normalization (Conv + BN + ReLU)

In practice, to increase the training as well as testing speed, we usually feed **multiple** images to the network. The following figure shows a training batch of 4 images,

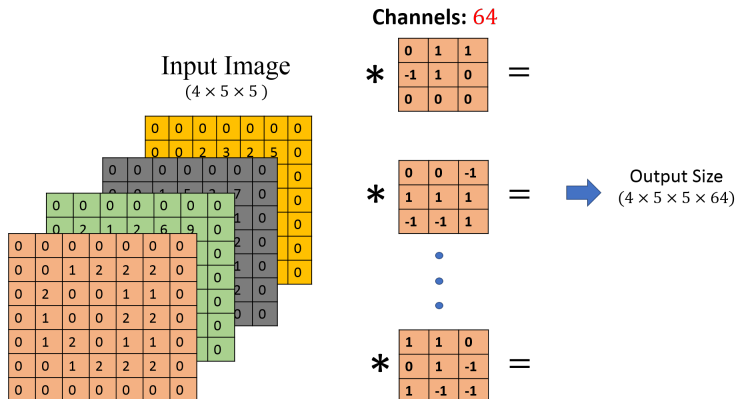


Figure: Batch Size is 4. Each Image is independently processed.

Batch Normalization : (Conv + **BN** + ReLU)

For each *channel*, normalize the layers. Mean and variance are computed across all the values in each channel.

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_1 \dots x_m\}$;

Parameters to be learned: γ, β

Output: $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

Algorithm 1: Batch Normalizing Transform, applied to activation x over a mini-batch.


An effective way to resolve *vanishing gradient* problem!

Sergey Ioffe, Christian Szegedy. *Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift*, NIPs 2015.


Nonlinear Activations: (Conv + BN + **ReLU**)

Popular nonlinearities used through all **but** last layer:


- ReLU: $\max(0, x)$.
- Leaky ReLU: $\max(0, x) + \gamma^2 \min(0, x)$
- Tanh: $\frac{\exp(x) - \exp(-x)}{\exp(x) + \exp(-x)}$
- Others ³: ELU, SELU, PReLU, Threshold *etc.*



(a) ReLU



(b) Leaky ReLU



(c) Tanh

³A good place to find all these is the document of deep learning software frameworks, eg. <http://pytorch.org/docs/master/nn.html>

Pooling




Figure: Illustration of *Max Pooling*. Here batch size 1 is omitted. ⁴

There is also average pooling which takes average rather than maximum.

⁴Image courtesy of CS231n: Convolutional Neural Networks for Visual Recognition. Stanford.

Skip Connection




Figure: Skip connection is simply concatenating two feature maps (along channel dimension). Central crop is performed if there is a mismatch of the dimension.

Comment. Adding skip connections can usually increase the performance (at least) in segmentation.

Now We Know $(+ - \times \div)$...


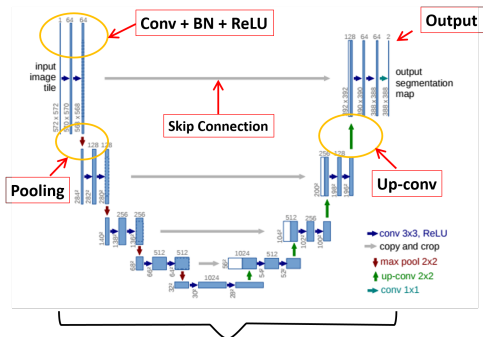


Figure: Networks are just special ways to stack all these operations.

Function Representation



$$f(I; W)$$

Figure: The entire network is just another representation of $f(I; W)$, where I is input image(s), W are the parameters in the network.

Optimization Problem

In training state, we are given ground truth images and its labels for recognition/classification, or masks for segmentation, high-resolution image for super-resolution, clear image for de-noising *etc.*.

Let I be the truth images and g be their labels, the optimization problem is


$$\min_W \text{Loss}\{f(I; W) - g\}$$

Popular losses,

- l_2 **distance**: image denoising, super-resolution, recognition.
- **Entropy**: binary/categorical cross-entropy, KL-divergence for segmentation.
- Many other **customized losses**, *e.g.* l_1 for GAN, smoothed dice loss for segmentation.

Loss Function Example: Cross Entropy

Suppose we are doing a single object segmentation, *e.g.* lung. The last convolutional channel will have 2 channels corresponding foreground and background.



Softmax + Cross-entropy

Outline

- ① Overview
- ② Convolutional Neural Networks Basics
 - Basic Operations
 - Optimization Problems in CNNs
- ③ Optimization
- ④ Networks Variants
- ⑤ Coding

Overview

Outline

- ① Overview
- ② Convolutional Neural Networks Basics
 - Basic Operations
 - Optimization Problems in CNNs
- ③ Optimization
- ④ Networks Variants
- ⑤ Coding

Networks

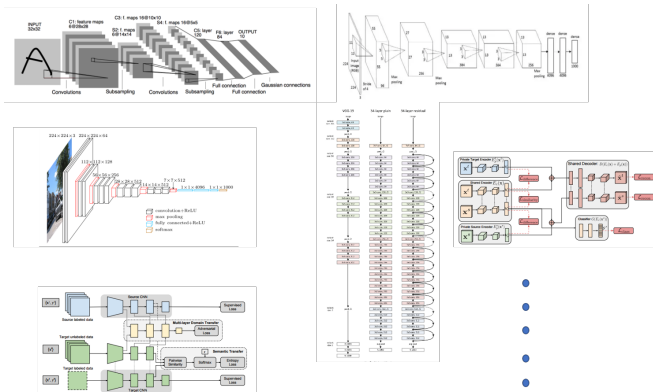


Figure: Networks are just special ways to stack all these operations.

Outline

- ① Overview
- ② Convolutional Neural Networks Basics
 - Basic Operations
 - Optimization Problems in CNNs
- ③ Optimization
- ④ Networks Variants
- ⑤ Coding

Frameworks

THANK YOU!